

RDBMS:OLAP & OLTP

-Aman Khandwe

March 25, 2022

Abstract

A relational database is a digital database that uses the relational data model. A relational database management system is a system for maintaining relational databases (RDBMS). Many relational database systems allow you to query and maintain the database using SQL (Structured Query Language). It's a programme that lets us create, delete, and update relational databases. A relational database is a database system that stores and retrieves data in the form of rows and columns in a tabular format. The main DBMS, such as SQL, My-SQL, and ORACLE, are all based on relational DBMS concepts. The fact that the values of each table are connected to each other is the cornerstone of relational DBMS. It is capable of handling enormous amounts of data and easily simulating queries. Here in this paper we will look at the relational model and will look at the terminologies related to RDBMS and how to store objects in relational database, followed by advantages, disadvantages, and some applications of the same, and finally we will try to have a deeper look at the two major applications of RDBMS, OLAP (Online analytical processing) and OLTP (Online transaction processing) (Online transaction processing).

Introduction

RDBMS is a concept that is used to give the user with the ability to store and retrieve data that is connected with any other relation, also known as the Relational. A relation in a relational database is based on a relational scheme and is made up of a number of properties. A relational database is composed of several relations and a relational database architecture. Users have the ability to access and alter the data stored in the database in a convenient and functional manner. Furthermore, the DBMS exerts centralised control over the database, prohibits unauthorised users from accessing the

data, and secures data privacy. The fact that the values of one table are connected to others is the foundation of relational DBMS. It is capable of handling bigger amounts of data and quickly simulating queries. Relational Database Management Systems ensure data integrity by imitating the following characteristics:

1. Entity Integrity: No two records in the database table may be identical.
2. Referential Integrity: Only those rows of those tables that are not utilised by other tables can be erased. Otherwise, data discrepancy may occur.
3. User-defined Integrity: Rules based on confidentiality and access that are established by the users.
4. Domain integrity: The database table columns are contained inside certain specified limitations based on default values, data type, or ranges.

In RDBMS, data must be saved in a DB file in tabular form, that is, in the form of rows and columns. Each table row is referred to as a record/tuple. The cardinality of the table refers to the collection of such items. Each column in the table is referred to as an attribute/field. The arity of the table is a collection of similar columns. There can be no duplicate records in the DB table. By employing a candidate key, data duplication is eliminated. A Candidate Key is a collection of properties that must be present in order for each record to be uniquely identified. Tables are linked to one another via foreign keys. Database tables also support NULL values, which means that if the values of any of the table's elements are not filled in or are missing, the value becomes a NULL value, which is not equal to zero. RDBMS often include data dictionaries and metadata collections that aid in data management. These provide for the programmatic support of well-defined data structures and relationships. Data storage management is a typical RDBMS functionality, and it has been characterised by data objects ranging from binary large object strings to stored procedures. This type of data item extends the scope of conventional relational database operations and may be handled in a number of ways by different RDBMSes. For accessing data in RDBMS SQL is used. Data manipulation language and data definition language statements are its primary language components. RDBMS employs complicated algorithms that allow several concurrent users to access the database while preserving data integrity. Another overlay function provided by the RDBMS for the fundamental database when used in business settings is security management, which enforces policy-based access.

Relational Model

The relational model is based on the mathematical idea of a relation, which is represented practically as a table. Relations are utilized in the relational model to store information

about the items that will be represented in the database. A relation is represented as a two-dimensional table, with rows corresponding to individual records and columns corresponding to attributes. Attributes can occur in any sequence, and the connection remains the same and so conveys the same meaning. Domains are a very important aspect of the relational model. A domain is used to specify each attribute in a relation. Domains might be unique for each characteristic, or they can be shared by two or more attributes. The domain idea is significant because it allows the user to specify the meaning and source of values that attributes can carry in a centralized location. As a result, the system has more information accessible to it when performing a relational action, and operations that are semantically inaccurate can be avoided. The rows or tuples in the table are the elements of a relation. Tuples can be arranged in any sequence, and the connection will remain the same and hence communicate the same meaning. The structure of a relation, together with a definition of the domains and any other constraints on potential values, is referred to as its intension, which is normally fixed unless the meaning of the relation is updated to incorporate more characteristics. The tuples are referred to as a relation's extension (or state), which varies over time. A unary relation, also known as a one-tuple, is a relation with only one attribute and has degree one. Binary refers to a relationship with two qualities, ternary refers to a relationship with three attributes, and n-ary refers to a relationship with more than three attributes. A relation's degree is a characteristic of the relation's intension. The cardinality of the relation, refers to the number of tuples in the relation, which changes when tuples are added or removed. The cardinality of a connection is a feature of its extension that is defined by the specific instance of the relation at any given time.

A relation has the following properties:

- The relation has a unique name that differs from the names of all other relations in the relational schema.
- There is only one atomic (single) value in each cell of the relation.
- Each attribute is given a unique name.
- An attribute's values are all from the same domain.
- There are no duplicate tuples; each tuple is unique.
- The order of the attributes is irrelevant.
- The order of tuples has no theoretical importance.

Integrity Constraints

A data model contains two additional components: a manipulative part that defines the sorts of operations that may be performed on the data, and a set of integrity constraints that verify the data's accuracy. The relational integrity restrictions will be discussed in this section. Since each attribute has a domain, there are constraints (known as domain constraints) that limit the set of values that may be assigned to the attributes of relations. There are also two key integrity rules, which are limits or limitations that apply to all database instances. Entity integrity and referential integrity are the two main rules of the relational model. Multiplicity and general constraints are two further forms of integrity constraints. It's important to grasp the idea of nulls before we can establish entity and referential integrity. Null denotes the absence of a value for an attribute that is presently unknown or irrelevant to this tuple. A null is not the same as a numeric value of zero or a text string with spaces; all are values, but a null signifies the lack of a value. So, nulls must be handled differently from other values. Now let us talk about the relational integrity rules.

Entity Integrity The main keys of base relations are subject to the first integrity rule. No attribute of a main key can be null in a base relation. A primary key, by definition, is a single identifier that is used to uniquely identify tuples. This indicates that no subset of the main key is adequate to identify tuples uniquely. Allowing a null for any component of a primary key implies that not all attributes are required to differentiate between tuples, which is contrary to the primary key's definition.

Referential Integrity Foreign keys are subject to the second integrity rule. If a relation has a foreign key, the value of the foreign key must either match a candidate key value of some tuple in the home relation or be completely null. Referential integrity constraints are used to preserve consistency between tuples in two relations and are expressed between two relations or tables.

General Constraints These are the additional rules set by database users or database administrators that define or limit a particular area of the business. Additionally, users can set additional constraints that the data must meet. For example, if the number of people who can work at a branch office is limited to 20, the user must be able to define this general limitation and expect the DBMS to enforce it. If the number of staff presently allocated to a branch is 20, it should not be able to add a new member of staff to the Staff relation in this situation. However, the amount of support for generic restrictions varies considerably from one system to the next.

Storing Objects in a Relational Database

Using an RDBMS as the underlying storage engine is one way to achieve persistence with an object-oriented programming language like C++ or Java. This necessitates mapping class instances (i.e., objects) to one or more tuples distributed across one or more relations. Consider the inheritance structure depicted in Figure 1, which includes a superclass called Staff and three subclasses called Manager, SalesPersonnel, and Secretary.

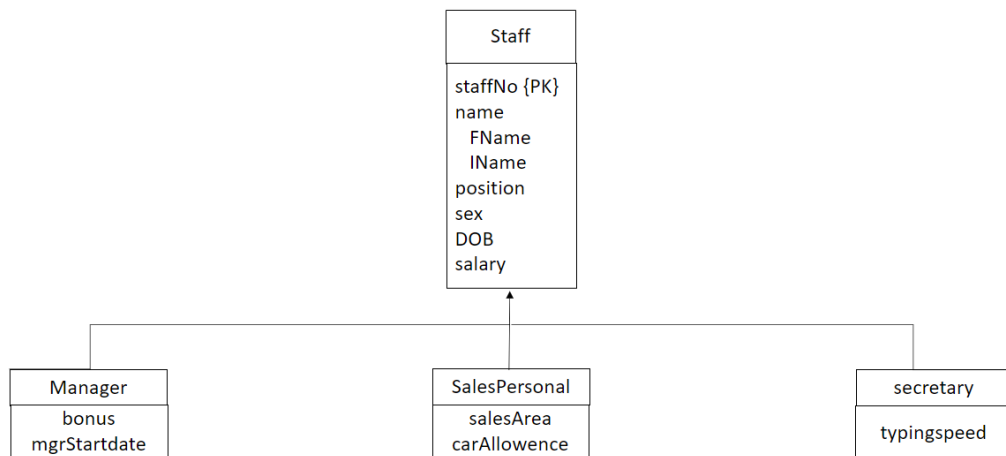


Figure 1: Inheritance hierarchy for Staff.

We have two main activities to undertake when dealing with this sort of class hierarchy:

1. Design the relations to represent the class hierarchy.
2. Create a plan for how objects will be accessible, which includes:
 - developing code that decomposes objects into tuples and stores them in relations;
 - writing code that reads tuples from relations and reconstructs the objects.

Now let us look at in depth about the above two mentioned tasks-

1.Mapping Classes to Relations

There are several strategies for mapping classes to relations, each of which results in a loss of semantic information. The code used to make objects persistent and read them back from the database is determined by the strategy used. We consider the following three options:

1. Map each class or subclass to a relation.
2. Map each subclass to a relation.
3. Map the hierarchy to a single relation.

1.1 Map each class or subclass to a relation

One method is to map each class or subclass with a relation. This would result in the following four relationships for the hierarchy shown in Figure 1:

1. Staff (staffNo, fName, lName, position, sex, DOB, salary).
2. Manager (staffNo, bonus, mgrStartDate).
3. SalesPersonnel (staffNo, salesArea, carAllowance).
4. Secretary (staffNo, typingSpeed).

Although we presume that the RDBMS supports the underlying data type of each attribute, this may not be the case, in that case we will need to create additional code to manage the translation from one data type to another. Unfortunately, we have lost semantic information with this relational schema: it is no longer evident which relation represents the superclass and which relation represents the subclasses. As a result, we'd have to embed this knowledge into each programme, which, as we've discussed before, can lead to code duplication and inconsistency.

1.2 Map each subclass to a relation

A second option is to map each subclass with a relation. This would result in the following three relations for the hierarchy shown in Figure 1:

1. Manager (staffNo, fName, lName, position, sex, DOB, salary, bonus, mgrStartDate)

2. SalesPersonnel (staffNo, fName, lName, position, sex, DOB, salary, salesArea, car-Allowance)
3. Secretary (staffNo, fName, lName, position, sex, DOB, salary, typingSpeed)

Again, semantic information has been lost in this mapping: it is no longer evident whether these relations are subclasses of a single generic class. To get a list of all employees in this example, we'd have to choose the tuples from each relation and then combine the results.

1.3 Map the hierarchy to a single relation

A third option is to translate the complete inheritance hierarchy to a single relation, as shown below:

Staff (staffNo, fName, lName, position, sex, DOB, salary, bonus, mgrStartDate, salesArea, carAllowance, typingSpeed, typeFlag)

typeFlag is a discriminator that determines which type each tuple belongs to (for example, it may contain the value 1 for a Manager tuple, 2 for a SalesPersonnel tuple, and 3 for a Secretary tuple). In this mapping also we've lost semantic information. This mapping will also result in a large number of nulls for attributes that do not apply to that tuple. The properties salesArea, carAllowance, and typingSpeed, will be null for a Manager tuple.

Advantages of RDBMs

A Relational Database system has numerous advantages over other types of databases. The following are a few important advantages:

1.Simple Model

A Relational Database system is the simplest basic type since it does not need any sophisticated structure or querying methods. It does not include time-consuming architecture operations such as hierarchical database structuring or definition. Because the structure is basic, it can be handled with simple SQL queries and does not necessitate the creation of sophisticated queries.

2.Data Accuracy

Multiple tables in a relational database system can be linked to one another using the primary key and foreign key principles. As a result, the data is non-repetitive. There is no possibility of data duplication. As a result, the accuracy of data in a relational database is greater than that of any other database system.

3.Easy Access to Data

There is no pattern or method for accessing data in the Relational Database System, whereas other types of databases can only be accessed by travelling via a tree or a hierarchical architecture. Anyone with data access can query any table in the relational database. Using join queries and conditional expressions, you may combine all or any number of linked tables to get the data you need. The resulting data may be updated based on any column's values, on any number of columns, allowing the user to easily recover the relevant data as a consequence. It enables the user to select the desired columns to be included in the output, ensuring that only relevant data is presented.

4.Data Integrity

The integrity of data is a critical feature of the Relational Database system. Sturdy Data entries and validity validations guarantee that all data in the database is contained within appropriate arrangements and that the data required to create relationships is available. This relational consistency throughout the database tables helps to prevent records from being incomplete, isolated, or unconnected. Data integrity helps to ensure the relational database's other important features, such as ease of use, accuracy, and data stability.

5.Flexibility

A Relational Database system, on its own, contains attributes for scaling up and extending for greater lengths, since it is supplied with a malleable structure to fit continually changing requirements. This makes it easier to deal with the rising volume of data coming in, as well as to update and delete records as needed. This model accepts modifications made to a database setup as well, which may be done without crashing the data or other elements of the database. Individuals can quickly and simply insert, amend, or remove tables, columns, or individual data in the specified database system. A relational database is said to have no limit on the amount of rows, columns, or tables it may contain. The Relational Database Management System and the hardware provided

by the servers limit development and transformation in any real application. As a result, these alterations may cause changes in various peripheral functional devices linked to the specific relational database system.

6.Normalization

The systematic approach is maintained to ensure that a relational database structure is free of any variations that might affect the integrity and correctness of the database's tables. A normalisation process establishes a set of rules, features, and goals for a relational database model's database structure and assessment. The goal of normalisation is to show several levels of data breakdown. Before going on to the following levels of normalisation, any level of normalisation is anticipated to be completed on the same level. Only after a relational database model meets the requisite requirements of the third normalisation form is it typically determined to be normalised. Normalization gives the idea that the database design is particularly sturdy and dependable.

7.High Security

Since the data is distributed across the tables of the relational database system, a few tables might be marked as confidential while others are not. Unlike other databases, a relational database management system makes this separation simple to build. When an individual attempts to log in with a username and password, the database can limit their degree of access by allowing them to access just the tables that they are authorised to work on, based on their access level.

8.Feasible for Future Modifications

Since the relational database system organises entries into tables depending on their categories, it's simple to add, remove, or update records that meet the most recent needs. This aspect of the relational database model allows the business to accommodate new requirements. By adhering to the fundamental properties of the relational database management system, any number of new or existing tables or columns of data may be introduced or updated based on the requirements supplied.

Disadvantages of RDBMs

1. Cost

The initial investment in a relational database is extremely high. A separate piece of software must be acquired in order to set up a relational database. In addition, the system should be maintained by a trained technician. All of this may be expensive, especially for small enterprises. Because relational databases have numerous appealing characteristics, they are rather expensive to utilise. As a result, obtaining such a database may seem difficult if the company has a limited budget.

2. Performance Issue

If there are a large number of rows and tables, the query will take longer to process the result because relational databases rely on rows and columns, performance might be sluggish. Furthermore, if there is a large amount of data in the system, it might slow down the working process, the presence of extra data not only slows down the system but also makes it more difficult to discover information. As a result, a relational database is recognised as a slower database. So we can say that the performance problem is a drawback of a relational database that most users may encounter when using it.

3. Physical Storage and Information Loss

Because it is made up of rows and columns, a relational database needs a large amount of physical memory. Each action is dependent on its own physical storage. Only by properly optimising the intended programmes can they be made to have the most physical memory. The RDBMS has limited space, and these storage devices cannot store new data if there is insufficient space. Because there is no longer any storage available, this data may be lost, causing issues in the future. As a result, relational databases are restricted and can cause a variety of problems if they are not adequately monitored.

4. Structure Limitations

A relational database has fields that have limits. In essence, limitations mean that it cannot handle additional information. It is vital to describe the data volume you wish to incorporate into any field while building the database. Because some of the search queries are or may be more exact than the original ones, data loss may occur. Even if

additional information is supplied, data loss may occur. As a result, the exact amount of data volume that the field will be given must be defined.

Applications of RDBMs

In the recent decade, there have been tremendous developments in the computer sector. RDBMSs have gained significant popularity in database systems for classic commercial applications such as order processing, inventory control, banking, and airline bookings. However, standard RDBMSs have proven insufficient for applications with requirements that differ significantly from those of traditional corporate database applications. These applications include:

- computer-aided manufacturing (CAM)- In addition to data pertaining to discrete production (such as automobiles on an assembly line), a CAM database includes data relevant to continuous production (such as chemical synthesis). There are applications that monitor information about the condition of the system, such as reactor vessel temperatures, flow rates, and yields, in chemical manufacture. There are additional applications that regulate various physical processes, such as opening valves, boosting the flow of cooling systems, and applying more heat to reactor vessels. These applications are frequently structured in a hierarchical structure, with a top-level application overseeing the whole factory and lower-level apps overseeing specific production processes.
- computer-aided software engineering (CASE)- The stages of the software development lifecycle are stored in a CASE database: planning, requirements collecting and analysis, design, implementation, testing, maintenance, and documentation. Designs can be exceedingly huge, much like in CAD, and cooperative engineering is the norm. Software configuration management technologies, for example, enable simultaneous sharing of project design, code, and documentation. They also help with change management by tracking the dependencies between various components. Project management tools make it easier to coordinate a variety of project management duties, such as scheduling potentially extremely complicated interdependent tasks, estimating costs, and tracking progress.